

ii. **Correctness Proof:** I will show that the equality between original problem and single source shortest path problem in G' :

- because for every v_s , it has only one out-edge points to v_e , if we enter v_s , we must pass (v_s, v_e) and pay $w(v_s, v_e) = c(v)$.
- if we are at u_e , we must have passed u_s and pay $c(u)$. Then, we can go to v_s if we can reach v from u in the original graph, which means in the next step we must pass (v_s, v_e) and pay $c(v)$ to "visit" city v .

Thus, the 2 problem is exactly equivalent, so the algorithm is correct.

Problem E - Mailing Fees

- (1) This problem can be modeled as a **Single Source Shortest Path** problem, where node 0 is the source vertex. For the "sending back" part, because the graph is undirected, so the shortest path from source to a vertex is equivalent to the shortest path from the vertex to source. Hence, I can run *Dijkstra's Algorithm* first to find the shortest distance to each factory and make them double, then they are the cost of a "round trip", which is the answer of the problem.

Mailing Fees I: Dijkstra

```

function Dijkstra(G, w):
    src ← 0
    // initializing
    for i from 1 to N - 1:
        dist[i] ← ∞
    dist[src] ← 0
    // initialize priority queue of 2-tuple,
    // which is ordered by the first element (distance)
    // with the minimum value on the top
    pq ← ∅
    INSERT(pq, (0, src))

    while pq is not empty:
        d, v ← EXTRACT.MIN(pq)
        // test if this element should be ignored
        if d ≠ dist[v]:
            continue
        // relax each edge
        for u in G.adj[v]:
            if d + w(v, u) < dist[u]:
                dist[u] ← d + w(v, u)
                INSERT(pq, (dist[u], u))

    // return the array that contains the shortest path to all vertices
    return dist

function solve(G, w):
    dist ← Dijkstra(G, w)
    for i from 0 to N - 1:
        dist[i] ← dist[i] * 2
    return dist

ANSWER ← solve(G, w)

```

- (2) Now the edge becomes directed, so a path from source to a vertex does not equal to a path from the vertex to source. To find the shortest path from each vertex to the source, I run *Dijkstra's Algorithm* on G^T , i.e. the graph that reverses the direction of every edge, because in G^T , a path from the source to a vertex is equivalent to a path from the vertex to the source in G .

The time complexity of the above implementation of *Dijkstra's Algorithm* is analyze as follow:

- (a) initializing dist[]: $O(N)$
- (b) initializing priority queue(binary min-heap): $O(1)$
- (c) In every iteration of the while loop: $O(V \log E) + O(E \log E) = O(E \log E)$
 - the min-element is extracted from pq: $O(\log E)$
 - relax all edges leaves v , upon successfully relaxing, insert the better weight and node id to pq: $O(\text{out-degree}(v) * \log E)$

* Note that the upper bound of the size of pq is E if every time we relax an edge, we insert a new element to pq.

....The overall complexity is $O(E \log E)$

We run *Dijkstra's Algorithm* twice, one on G , one on G^T (can be constructed in $\Theta(E)$). And we need to add the 2 result together, which takes $O(N)$ trivially. Therefore, the overall time complexity of my algorithm is

$$\Theta(E) + 2 \times O(E \log E) + O(N) = O(E \log E)$$

- (3) Now the edges can have negative weights, so I turn to *Bellman-Ford's Algorithm*, while *Dijkstra* cannot handle this scenario. The implementation is as follow:

Mailing Fees II: Bellman-Ford

```

function BellmanFord(G, w):
    src ← 0
    // initializing
    for i from 1 to N - 1:
        dist[i] ← ∞
    dist[src] ← 0
    // N - 1 rounds
    for i from 1 to N - 1:
        for (u, v) in G.E:
            if dist[u] + w(u, v) < dist[v]:
                dist[v] = dist[u] + w(u, v)

    // detect negative cycle
    for (u, v) in G.E:
        if dist[u] + w(u, v) < dist[v]:
            return FALSE

    return dist

function solve(G, w):
    result ← BellmanFord(G, w)
    if result is FALSE:
        print("I am rich!")
        return -1
    else:
        return dist

ANSWER ← solve(G, w)

```

The time complexity is analyzed as follow:

- (a) initializing: $O(N)$
- (b) $|V| - 1$ iterations of relaxation of all edges: $O(N * E)$
- (c) detect negative cycle: $O(E)$

Thus, the overall time complexity is $O(N * E)$

Correctness: The ability to find the shortest path and detect negative cycle of *Bellman-Ford's Algorithm* is proved in the class.

- (4) Here is a comparison table:

* NOTE: Let M_G be the space complexity to store the graph with adjacency lists, $M_G = O(V + E)$

	subproblem (2)	subproblem (3)
Algorithm	Dijkstra	Bellman-Ford
Time Complexity	$O(E \log E)$	$O(VE)$
Space Complexity	$O(V + E) + M_G$	$O(V) + M_G$
Advantages	better time complexity	can handle negative weights, can detect negative cycle
Disadvantages	cannot handle negative weights	worse time complexity

- (5) • **Algorithm Description:** I reweight the edges such that a "trustful cycle" is equal to a "negative cycle" in the graph. Let graph be $\langle V, E \rangle$ the original weight function be w , and $R(v)$ denotes the reliability value of factory v , then the new weight function w' is defined as follow:

$$\forall (u, v) \in E, w'(u, v) = K \times w(u, v) - R(v)$$

Then I run *Bellman-Ford's algorithm* to take use of its capability of detecting negative cycle. If the return value is FALSE, which means there exists a negative cycle, then we know there exists a "trustful cycle".

- **Correctness Proof:** I show that a "trustful cycle" in G is equivalent to a "negative cycle" in G' :

$$\begin{aligned}
 & \text{Let } C \text{ be a "trustful cycle" in } G \\
 \implies & \sum_{v \in C} R(v) \div \sum_{(u,v) \in C} w(u,v) > K \\
 \implies & K \times \sum_{(u,v) \in C} w(u,v) - \sum_{v \in C} R(v) < 0 \\
 \implies & \sum_{(u,v) \in C} K \times w(u,v) - R(v) < 0 \\
 \implies & \sum_{(u,v) \in C} w'(u,v) < 0
 \end{aligned}$$

- **Time Complexity Analysis:**

(a) reweighting: $O(E)$

(b) Bellman-Ford: $O(V * E) = O(N * E)$

Thus, the overall complexity is $O(E) + O(N * E) = O(N * E)$ #

Problem F - Gaussian is Too Slow

I model this problem as a *Minimum Spanning Tree* problem. Construct a undirected graph $G = \langle V, E \rangle$ and the weight function w as follow:

$$\begin{aligned} V &= \{v_0, v_1, v_2, \dots, v_N\} \\ E &= \{(v_i, v_j) \mid 0 \leq i < j \leq N\} \\ w(v_i, v_j) &= \text{the cost of } x_{i+1} + x_{i+2} + \dots + x_j, 0 \leq i < j \leq N \end{aligned}$$

After constructing the graph, run *Prim's algorithm* on it, then the total weight of result minimum spanning tree is the answer. I will prove the correctness and time complexity of the algorithm in the following.

- **Proof of Correctness:** (For simplicity, let $f_{i,j} = x_i + x_{i+1} + \dots + x_j$ in the following)

Claim: If there exist a simple path from v_i to v_j , WLOG assume $i < j$, then we can obtain $f_{i+1,j} = x_{i+1} + x_{i+2} + \dots + x_j$ by adding or subtracting the equations corresponded to the edges in the path.

Proof of Claim: Let the number of edges in the simple path be k , prove by induction on k :

- **Base Case:** $k = 1$, trivial, we only have one edge and one equation, and the equation is just what we want.
- **Inductive Step:** Assume the claim is correct when $k = m$, then I show that it is also correct when $k = m + 1$. If we have a simple path from v_i to v_j that consists of $m + 1$ edges, then the path can be split into 2 parts, one from v_i to some other node v_x that has m edges and another from v_x to v_j has 1 edges. By induction hypothesis, we can obtain $f_{i,x}$ if $i < x$ or $f_{x,i}$ if $x < i$:
 - * $i < x$:
 - $j < x$: we have $f_{i+1,x}$ and $f_{j+1,x} \implies f_{i+1,x} - f_{j+1,x} = f_{i+1,j}$
 - $j > x$: we have $f_{i+1,x}$ and $f_{x+1,j} \implies f_{i+1,x} + f_{x+1,j} = f_{i+1,j}$
 - * $i > x$:
 - $j < x$: in this case $j < x < i$, which contradicts $i < j$, so it is impossible
 - $j > x$: we have $f_{x+1,i}$ and $f_{x+1,j} \implies f_{x+1,j} - f_{x+1,i} = f_{i+1,j}$

By induction, the claim is true for all $k \geq 1$.

By claim, if a set of edges $T \subseteq E$, such that for all $i \in \{1, \dots, N\}$, there exists a simple path from v_{i-1} to v_i (i.e. all nodes are connected), then we can obtain $\{x_1, x_2, \dots, x_N\}$ by adding or subtracting the equations corresponded to the edges, which means each variable is solvable. Furthermore, a minimum spanning tree of G makes G connected with the total weight of edges in tree minimized. Hence, an algorithm for *Minimum Spanning Tree* like *Prim's* outputs the optimal answer of the original problem.

- **Time Complexity Analysis:** Since in this case, $|E| = C_2^{N+1} = \frac{N(N+1)}{2} = \Theta(N^2) \gg |V| = N + 1 = \Theta(N)$, we implement the priority queue in *Prim's algorithm* by an array indexed by the id of nodes, such that the operations have the following complexity:
 - extract-min: $O(|V|)$
 - insert: $O(1)$
 - decrease-key: $O(1)$

Then *Prim's algorithm* run in $|V| \times O(|V|) + O(|E|) \times O(1) = O(|E|) = O(N^2)$ time complexity.